



Different than the rest of machine learning

Predictive machine learning



Reinforcement learning



- Markov Decision Processes (MDP)
- The Bellman equation
- Q-networks
- Policy gradients

Temporal Difference (TD) Learning



Learn at each time step



Each network gets a cost

$$\delta = R_t + \gamma V(S_{t+1}) - V(S_t)$$

$$\delta^2$$



Critic loss

$$\delta \ln \pi(A_t | S_t)$$

Actor loss

**Values, states,
actions, and
policy**

Value and policy



Bellman equation

Value of state as a maximum of the value of its neighbors for the neighbors states obtained by applying any possible action

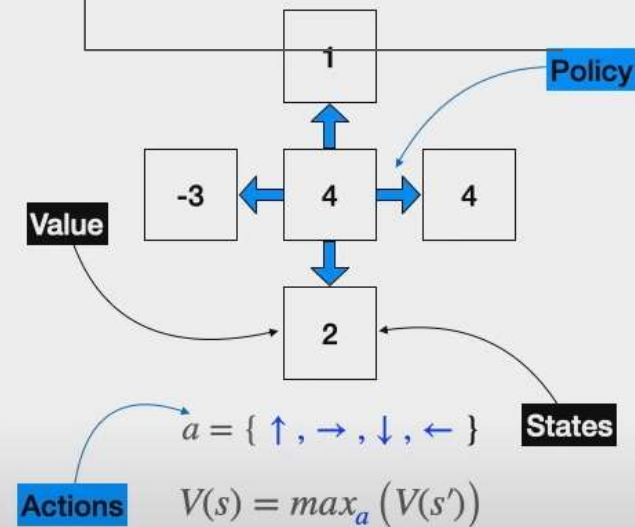

Best possible decision

\Leftrightarrow POLICY

POLICY (best decision)

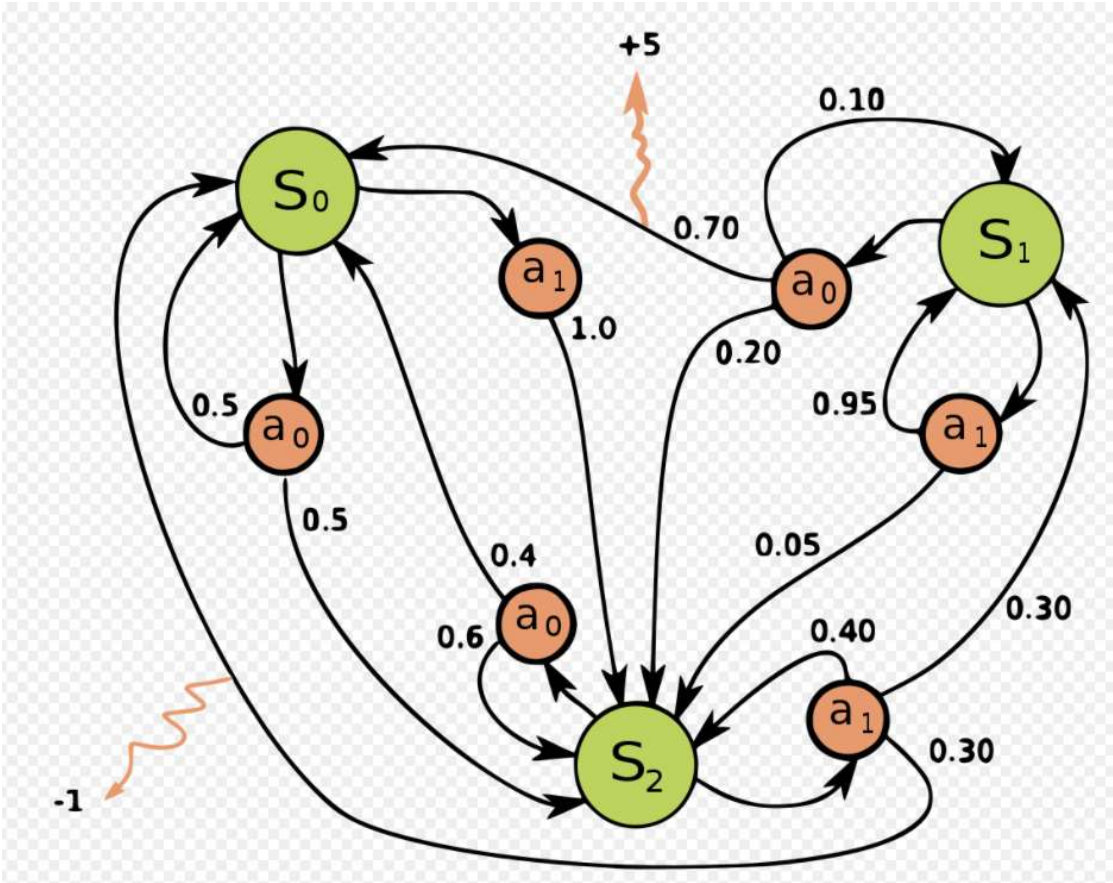
-> Set of instructions to always take the agent in the best possible path with respect to gaining points

Value and policy



Markov decision processes (MDP)

Example of a simple MDP with three states (green circles) and two actions (orange circles), with two rewards (orange arrows).





REWARD



-> Positive

-> Negative

Reward



Reward





$$V(s) = \max_a (R(s, a) + V(s'))$$

$$a = \{ \uparrow, \rightarrow, \downarrow, \leftarrow \}$$

Bellman equation 1st change


each value of state is the maximum of points you can obtain if you walk in the best possible way, & that best possible way is given by policy (point to the neighbor that has the highest value notice)

Reward



Discount factor



Discount factor



Bellman equation 2nd change

Value of state as a maximum over all its neighbors where all the states i can obtain by applying the actions of a DISCOUNT FACTOR γ times the value of that new state

Discount factor




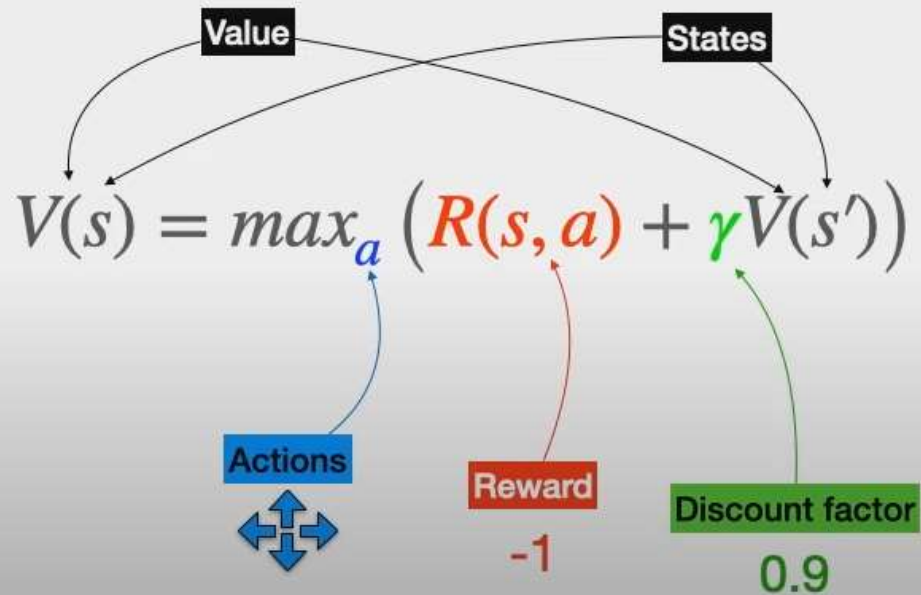
$$V(s) = \max_a (\gamma V(s'))$$

$$a = \{ \uparrow, \rightarrow, \downarrow, \leftarrow \}$$

Bellman equation

Bellman equation

$$V(s) = \max_a (R(s, a) + V(s')) \quad V(s) = \max_a (\gamma V(s'))$$



Solving the Bellman equation

| | | | | | |
|----|---|----|----|----|----|
| 2 | 1 | 2 | 3 | 4 | +5 |
| 3 | 2 | 1 | 2 | 3 | 4 |
| +4 | | 0 | 1 | 2 | 3 |
| 3 | | -1 | | | |
| 2 | 1 | 0 | -1 | -2 | -2 |
| 1 | 0 | -1 | -2 | -2 | -1 |

| | | | | | |
|------|-------|------|-------|------|------|
| 2 | ← 1 → | 2 | → 3 → | 4 | → +5 |
| ↓ 3 | ← 2 | 1 | → 2 | → 3 | → 4 |
| ↓ +4 | | 0 | → 1 | → 2 | → 3 |
| 3 | | ↑ -1 | | | |
| ↑ 2 | ← 1 | ← 0 | ← -1 | ← -2 | ← -2 |
| ↑ 1 | ← 0 | ← -1 | ← -2 | ← -2 | ← -1 |


Bellman equation allways satisfied

Best possible decision
⇔ POLICY

If we have a much bigger example : the algo says : you have to visit all the boxes, not only once but several times until the values start converging & that's impossible to do when you have a very very large universe with millions of states & actions... : it is very expensive to go over all of them, so what can we do ?


THIS IS WHERE NEURAL NETWORK COME TO OUR RESCUE

NN will come in 2 ways : one of them for calculating the value which is called a value network (Q NETWORK) & the other one for calculating the policy (Policy NETWORK)




In this way we won't have to go over every state several times to learn the value of the policy we simply have to let our agent wander and the NN will be smart enough to pick up information from the places that this agent manages to visit & propagate them to the entire grid.


This not only good for saving time but also for saving space



Deterministic vs stochastic policies




EXPLOITATION




EXPLORATION

Deterministic vs stochastic

Deterministic






Stochastic




Neural networks

Policy network




Two neural networks

Value neural network




Policy neural network



Value neural networks

Value neural networks



$$V(s) = \max_a (R(s, a) + \gamma V(s))$$

Bellman equation TO
BE always satisfied

| | | | |
|---|------|------|-----|
| 3 | | 1.3 | |
| 2 | -2.7 | 0.45 | 3.2 |
| 1 | | 4.9 | |
| | 2 | 3 | 4 |






Error function

$$\left(V(s) - \max_a \left(R(s, a) + \gamma V(s) \right) \right)^2$$


$$(0.2 - \max\{2.2, 0.3, -3.7, 3.9\})^2$$

Policy neural networks


Policy neural network




How to train it?




| Gain | x | y | Direction | Change |
|------|-----|-----|-----------|----------|
| 3 | 0 | 2 | ↑ | increase |
| 2 | 0 | 1 | ↑ | increase |
| 1 | 1 | 1 | ← | increase |
| 0 | 1 | 0 | ↑ | stay |
| -1 | 0 | 0 | → | decrease |





we take another walk



| Gain | x | y | Direction | Change |
|------|---|---|-----------|----------|
| 3 | 0 | 2 | ↑ | increase |
| 2 | 0 | 1 | ↑ | increase |
| 1 | 1 | 1 | ← | increase |
| 0 | 1 | 0 | ↑ | stay |
| -1 | 0 | 0 | → | decrease |



That is the table we give to NN

Training the policy neural network

Training the policy network



| Gain | x | y | action | Probability | Increase log loss | change |
|------|---|---|--------|------------------------|-------------------------|----------|
| 4 | 4 | 5 | → | $p(\rightarrow) = 0.3$ | 4 $\ln(p(\rightarrow))$ | increase |
| 3 | 4 | 4 | ↑ | $p(\uparrow) = 0.9$ | 3 $\ln(p(\uparrow))$ | increase |
| 2 | 3 | 4 | → | $p(\rightarrow) = 0.1$ | 2 $\ln(p(\rightarrow))$ | increase |
| 1 | 2 | 4 | → | $p(\rightarrow) = 0.2$ | 1 $\ln(p(\rightarrow))$ | increase |
| 0 | 2 | 3 | ↑ | $p(\uparrow) = 0.5$ | 0 $\ln(p(\uparrow))$ | stay |
| -1 | 2 | 2 | ↑ | $p(\uparrow) = 0.4$ | -1 $\ln(p(\uparrow))$ | decrease |
| -2 | 2 | 1 | ↑ | $p(\uparrow) = 0.3$ | -2 $\ln(p(\uparrow))$ | decrease |
| -3 | 3 | 1 | ← | $p(\leftarrow) = 0.7$ | -3 $\ln(p(\leftarrow))$ | decrease |